

## House Sales – Price Predictions

This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015.

**id** :a notation for a house

**date**: Date house was sold

**price**: Price is prediction target

**bedrooms**: Number of Bedrooms/House

**bathrooms**: Number of bathrooms/bedrooms

**sqft\_living**: square footage of the home

**sqft\_lot**: square footage of the lot

**floors** :Total floors (levels) in house

**waterfront** :House which has a view to a waterfront

**view**: Has been viewed

**condition** :How good the condition is Overall

**grade**: overall grade given to the housing unit, based on King County grading system

**sqft\_above** :square footage of house apart from basement

**sqft\_basement**: square footage of the basement

**yr\_built** :Built Year

**yr\_renovated** :Year when house was renovated

**zipcode**:zip code

**lat**: Latitude coordinate

**long**: Longitude coordinate

**sqft\_living15** :Living room area in 2015(implies-- some renovations) This might or might not have affected the lotsize area

**sqft\_lot15** :lotSize area in 2015(implies-- some renovations)

You will require the following libraries

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import
StandardScaler, PolynomialFeatures
%matplotlib inline
```

## 1.0 Importing the Data

Load the csv:

In [2]:

```
file_name='https://s3-api.us-geo.objectstorage.softlayer.net/cf-
courses-
data/CognitiveClass/DA0101EN/coursera/project/kc_house_data_NaN.
csv'
df=pd.read_csv(file_name)
```

we use the method head to display the first 5 columns of the dataframe.

In [3]:

```
df.head()
```

Out[3]:

	Unnamed: 0	id	date	price	bedrooms	bathrooms	sqft_lot_vin	sqft_lot	floors	waterfront	view	garage	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
0		7129300520	20141013T00:00:00	221900.0	3.0	1.0	1180	5650	1.0	0	0	0	1180	0	1955	0	98178	47.5112	-122.257	1340	5650
1		6414100192	20141020T00:00:00	538000.0	3.0	2.25	2570	7242	2.0	0	0	0	2170	400	1951	1991	98125	47.7210	-122.319	1690	7639

	Un na me d: 0	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
2	2	5631500400	20150225T00:00:00	180000.0	2.0	1.00	770	1000	1.0	0	.	6	770	0	1933	0	98028	47.7379	-122.233	2720	8062
3	3	2487200875	20141209T00:00:00	604000.0	4.0	3.00	1960	5000	1.0	0	.	7	1050	910	1965	0	98136	47.5208	-122.393	1360	5000
4	4	1954400510	20150218T00:00:00	510000.0	3.0	2.00	1680	8080	1.0	0	.	8	1680	0	1987	0	98074	47.6168	-122.045	1800	7503

5 rows × 22 columns

### Question 1

Display the data types of each column using the attribute dtype, then take a screenshot and submit it, include your code in the image.

In [5]:

```
print(df.dtypes)
Unnamed: 0      int64
id              int64
date            object
price           float64
bedrooms        float64
bathrooms       float64
sqft_living      int64
sqft_lot         int64
floors          float64
waterfront      int64
view            int64
condition       int64
grade           int64
sqft_above      int64
sqft_basement   int64
```

```

yr_built          int64
yr_renovated      int64
zipcode           int64
lat              float64
long             float64
sqft_living15     int64
sqft_lot15        int64
dtype: object

```

We use the method describe to obtain a statistical summary of the dataframe.

In [6]:

```
df.describe()
```

Out[6]:

	Unnamed: 0	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
count	2161	2130	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	2100	216000
mean	10806.0009	4.5803e+04	5.4881e+05	3.3720	2.1156	2079.899736	1.51069e+04	1.4943	0.07542	0.2343	7.656873	1788.390691	291.509045	1971.005136	84.402258	98077.939805	47.560053	-12.2213896	1986.552492	12768.455652
std	6239.28002	2.8765e+09	3.6712e+05	0.9267	0.768996	918.440897	4.142051e+04	0.5399	0.086517	0.766318	1.175459	828.090978	442.575043	29.373411	401.679240	53.505026	0.138564	0.140828	685.391304	27304.179631
min	0.0000	1.0010e+06	7.5000e+04	1.0000	0.5000	2900.0000	5.2000e+02	1.0000	0.0000	0.0000	1.0000	2900.0000	0.0000	1900.0000	0.0000	9800.0000	47.155900	-12.251900	399.0000	651.000000

	Un na me d: 0	id	pri ce	be dr oo ms	bat hr oo ms	sqf t_li vin g	sqf t_lo t	flo ors	wa ter fro nt	vie w	.	gra de	sqf t_a bo ve	sqf t_b ase me nt	yr_ bui lt	yr_ ren ov ate d	zip co de	lat	lon g	sqf t_li vin g1 5	sqft _lot 15
2	5403.000000	2.123049	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000	0.000000	.	7.000000	1190.000000	0.000000	1951.000000	0.000000	98033.000000	47.471000	-12.232800	1490.000000	510.000000
5	10806.000000	3.904930	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	.	7.000000	1560.000000	0.000000	1975.000000	0.000000	98065.000000	47.571800	-12.223000	1840.000000	762.000000
7	16209.000000	7.308900	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000	0.000000	.	8.000000	2210.000000	560.000000	1997.000000	0.000000	98118.000000	47.678000	-12.212500	2360.000000	10083.000000
m a x	21612.000000	9.900000	7.700000e+06	33.000000	8.000000	1354.000000	1.651359e+06	3.500000	1.000000	4.000000	.	13.000000	9410.000000	482.000000	2015.000000	2015.000000	98199.000000	47.777600	-12.131500	6210.000000	871200.000000

8 rows × 21 columns

## 2.0 Data Wrangling

### Question 2

Drop the columns "id" and "Unnamed: 0" from axis 1 using the method `drop()`, then use the method `describe()` to obtain a statistical summary of the data. Take a screenshot and submit it, make sure the `inplace` parameter is set to `True`

In [7]:

```
df.drop(['id', 'Unnamed: 0'], axis=1, inplace=True)
df.describe()
```

Out[7]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	sqft_living15	sqft_lot15
count	2.161300e+04	216000000	216030000	216130000	2.161300e+04	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000	216130000
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04	1.494309	0.007542	0.234303	3.409430	7.656873	1788.390691	291.509045	1971.005136	84.402258	98077.939805	47.560053	-122.213896	1986.552492	12768.455652
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04	0.539989	0.086517	0.766318	0.650743	1.175459	828.090978	442.575043	29.373411	401.679240	53.505026	0.138564	0.140828	685.391304	27304.179631
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02	1.000000	0.000000	0.000000	1.000000	1.000000	290.000000	0.000000	190.000000	0.000000	98001.000000	47.155900	-122.519000	399.000000	651.000000
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	0.000000	0.000000	3.000000	7.000000	1190.000000	0.000000	1951.000000	0.000000	98033.000000	47.471000	-122.328000	1490.000000	5100.000000
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	0.000000	0.000000	3.000000	7.000000	1560.000000	0.000000	1975.000000	0.000000	98065.000000	47.571800	-122.230000	1840.000000	7620.000000
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	0.000000	0.000000	4.000000	8.000000	2210.000000	560.000000	1997.000000	0.000000	98118.000000	47.678000	-122.125000	2360.000000	10083.000000
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	1.000000	4.000000	5.000000	13.000000	9410.000000	4820.000000	2015.000000	2015.000000	98199.000000	47.777600	-121.315000	6210.000000	871200.000000

we can see we have missing values for the columns bedrooms and bathrooms

In [8]:

```
print("number of NaN values for the column bedrooms :",  
      df['bedrooms'].isnull().sum())  
print("number of NaN values for the column bathrooms :",  
      df['bathrooms'].isnull().sum())  
number of NaN values for the column bedrooms : 13  
number of NaN values for the column bathrooms : 10
```

We can replace the missing values of the column 'bedrooms' with the mean of the column 'bedrooms' using the method replace. Don't forget to set the inplace parameter to True

In [9]:

```
mean=df['bedrooms'].mean()  
df['bedrooms'].replace(np.nan,mean, inplace=True)
```

We also replace the missing values of the column 'bathrooms' with the mean of the column 'bedrooms' using the method replace. Don't forget to set the inplace parameter to True

In [10]:

```
mean=df['bathrooms'].mean()  
df['bathrooms'].replace(np.nan,mean, inplace=True)
```

In [11]:

```
print("number of NaN values for the column bedrooms :",  
      df['bedrooms'].isnull().sum())  
print("number of NaN values for the column bathrooms :",  
      df['bathrooms'].isnull().sum())  
number of NaN values for the column bedrooms : 0  
number of NaN values for the column bathrooms : 0
```

### 3.0 Exploratory data analysis

#### Question 3

Use the method value\_counts to count the number of houses with unique floor values, use the method .to\_frame() to convert it to a dataframe.

In [12]:

```
df['floors'].value_counts().to_frame()
```

Out[12]:

	<b>floors</b>
<b>1.0</b>	10680
<b>2.0</b>	8241

	<b>floors</b>
<b>1.5</b>	1910
<b>3.0</b>	613
<b>2.5</b>	161
<b>3.5</b>	8

#### Question 4

Use the function `boxplot` in the seaborn library to determine whether houses with a waterfront view or without a waterfront view have more price outliers .

In [14]:

```
sns.boxplot(x='waterfront', y='price', data=df)
/opt/conda/envs/DSX-Python35/lib/python3.5/site-
packages/seaborn/categorical.py:462: FutureWarning: remove_na is
deprecated and is a private function. Do not use.
```

```
    box_data = remove_na(group_data)
```

Out[14]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f24f8d0b208>
```

#### Question 5

Use the function `regplot` in the seaborn library to determine if the feature `sqft_above` is negatively or positively correlated with price.

In [15]:

```
sns.regplot(x='sqft_above', y='price', data=df)
```

Out[15]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f24f8ca7d30>
```

We can use the Pandas method `corr()` to find the feature other than price that is most correlated with price.

In [16]:

```
df.corr()['price'].sort_values()
```

Out[16]:

```
zipcode          -0.053203
long             0.021626
condition        0.036362
yr_built         0.054012
sqft_lot15       0.082447
sqft_lot         0.089661
yr_renovated     0.126434
```



```
floors          0.256794
waterfront      0.266369
lat             0.307003
bedrooms        0.308797
sqft_basement   0.323816
view            0.397293
bathrooms       0.525738
sqft_living15   0.585379
sqft_above      0.605567
grade           0.667434
sqft_living     0.702035
price           1.000000
Name: price, dtype: float64
```

## Module 4: Model Development

### Import libraries

In [17]:

```
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
```

We can Fit a linear regression model using the longitude feature 'long' and calculate the  $R^2$ .

In [18]:

```
X = df[['long']]
Y = df['price']
lm = LinearRegression()
lm
lm.fit(X, Y)
lm.score(X, Y)
```

Out[18]:

```
0.00046769430149007363
```

### Question 6

Fit a linear regression model to predict the 'price' using the feature 'sqft\_living' then calculate the  $R^2$ . Take a screenshot of your code and the value of the  $R^2$ .

In [20]:

```
X = df[['sqft_living']]
Y = df['price']
lm = LinearRegression()
lm.fit(X, Y)
lm.score(X, Y)
```

Out[20]:

```
0.49285321790379316
```

## Question 7

Fit a linear regression model to predict the 'price' using the list of features:

In [21]:

```
features = ["floors", "waterfront","lat" ,"bedrooms"
,"sqft_basement" ,"view"
,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
```

the calculate the  $R^2$ . Take a screenshot of your code

In [22]:

```
X = df[features]
Y= df['price']
lm = LinearRegression()
lm.fit(X, Y)
lm.score(X, Y)
```

Out[22]:

```
0.65769516660374938
```

## this will help with Question 8

Create a list of tuples, the first element in the tuple contains the name of the estimator:

```
'scale'
```

```
'polynomial'
```

```
'model'
```

The second element in the tuple contains the model constructor

```
StandardScaler()
```

```
PolynomialFeatures(include_bias=False)
```

```
LinearRegression()
```

In [23]:

```
Input=[('scale',StandardScaler()),('polynomial',
PolynomialFeatures(include_bias=False)),('model',LinearRegression())]
```

## Question 8

Use the list to create a pipeline object, predict the 'price', fit the object using the features in the list features , then fit the model and calculate the R<sup>2</sup>

In [24]:

```
pipe=Pipeline(Input)
pipe
```

Out[24]:

```
Pipeline(memory=None,
          steps=[('scale', StandardScaler(copy=True, with_mean=True,
with_std=True)), ('polynomial', PolynomialFeatures(degree=2,
include_bias=False, interaction_only=False)), ('model',
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,
normalize=False))])
```

In [25]:

```
pipe.fit(X,Y)
```

Out[25]:

```
Pipeline(memory=None,
          steps=[('scale', StandardScaler(copy=True, with_mean=True,
with_std=True)), ('polynomial', PolynomialFeatures(degree=2,
include_bias=False, interaction_only=False)), ('model',
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,
normalize=False))])
```

In [26]:

```
pipe.score(X,Y)
```

Out[26]:

```
0.75134126473712171
```

## Module 5: MODEL EVALUATION AND REFINEMENT

import the necessary modules

In [27]:

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import train_test_split
print("done")
done
```

we will split the data into training and testing set

In [28]:

```
features =["floors", "waterfront","lat" ,"bedrooms"
,"sqft_basement" ,"view"
,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
X = df[features ]
Y = df['price']
```

```
x_train, x_test, y_train, y_test = train_test_split(X, Y,
test_size=0.15, random_state=1)
```

```
print("number of test samples :", x_test.shape[0])
print("number of training samples:",x_train.shape[0])
number of test samples : 3242
number of training samples: 18371
```

### Question 9

Create and fit a Ridge regression object using the training data, setting the regularization parameter to 0.1 and calculate the  $R^2$  using the test data.

In [29]:

```
from sklearn.linear_model import Ridge
```

In [30]:

```
RidgeModel = Ridge(alpha = 0.1)
RidgeModel.fit(x_train, y_train)
RidgeModel.score(x_test, y_test)
```

Out[30]:

```
0.64787591639391107
```

### Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, setting the regularisation parameter to 0.1. Calculate the  $R^2$  utilising the test data provided. Take a screenshot of your code and the  $R^2$ .

In [34]:

```
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import Ridge
pr = PolynomialFeatures(degree=2)
x_train_pr = pr.fit_transform(x_train)
x_test_pr = pr.fit_transform(x_test)
poly = Ridge(alpha=0.1)
poly.fit(x_train_pr, y_train)
poly.score(x_test_pr, y_test)
```

Out[34]:

```
0.70027442436889054
```

Once you complete your notebook you will have to share it. Select the icon on the top right marked in red in the image below, a dialogue box should open, select the option all content excluding sensitive code cells.